

```

def mean_squared_error(y_true, y_pred):
    return np.mean((y_true - y_pred) ** 2)

def mean_absolute_error(y_true, y_pred):
    return np.mean(np.abs(y_true - y_pred))

import numpy as np
y_true = np.array([1, 0, 1, 0])
y_pred = np.array([0.9, 0.1, 0.8, 0.3])
print("MSE:", mean_squared_error(y_true, y_pred))
print("MAE:", mean_absolute_error(y_true, y_pred))

MSE: 0.03749999999999999
MAE: 0.175

def binary_cross_entropy(y_true, y_pred):
    # Clip predictions to prevent log(0)
    y_pred = np.clip(y_pred, 1e-15, 1 - 1e-15)
    return -np.mean(y_true * np.log(y_pred) + (1 - y_true) * np.log(1 - y_pred))

print("Binary Cross-Entropy:", binary_cross_entropy(y_true, y_pred))

Binary Cross-Entropy: 17.26978799617044

class LinearRegression:
    def __init__(self, learning_rate=0.01, n_iterations=1000):
        self.learning_rate = learning_rate
        self.n_iterations = n_iterations
        self.theta = None

    def fit(self, X, y):
        m = len(y)
        X_b = np.c_[np.ones((m, 1)), X] # Add bias term
        self.theta = np.random.randn(2, 1) # Random initialization

        for iteration in range(self.n_iterations):
            y_pred = X_b.dot(self.theta)
            gradients = 2/m * X_b.T.dot(y_pred - y)
            self.theta -= self.learning_rate * gradients

    def predict(self, X):
        X_b = np.c_[np.ones((len(X), 1)), X] # Add bias term
        return X_b.dot(self.theta)

# Generate synthetic data for testing
X = 2 * np.random.rand(100, 1)
y = 4 + 3 * X + np.random.randn(100, 1)
# Train the model
model = LinearRegression(learning_rate=0.1, n_iterations=1000)
model.fit(X, y)

```

```
y_pred = model.predict(X)
```

```
y_pred
```

```
array([[6.95867313],  
       [5.55784819],  
       [5.03268131],  
       [5.42344461],  
       [5.88218218],  
       [5.69996301],  
       [8.69996135],  
       [9.25016529],  
       [5.50059219],  
       [9.2855695 ],  
       [4.51937017],  
       [9.20333699],  
       [6.59185506],  
       [5.58589989],  
       [6.77086405],  
       [7.42032303],  
       [4.05445206],  
       [6.7545722 ],  
       [7.90817007],  
       [5.04426698],  
       [5.42355594],  
       [5.77347619],  
       [8.62922954],  
       [4.57187918],  
       [7.18280602],  
       [4.74174123],  
       [8.85884954],  
       [8.89030681],  
       [5.18663491],  
       [4.16995045],  
       [4.82405247],  
       [8.93688952],  
       [9.53872983],  
       [9.78630927],  
       [7.7527572 ],  
       [9.81695496],  
       [8.33365688],  
       [9.76516085],  
       [7.07632048],  
       [8.79123097],  
       [7.4664544 ],  
       [8.92197924],  
       [7.34543062],  
       [4.9018413 ],  
       [9.08308922],  
       [6.21478645],
```

[8.74984961],  
[5.31660092],  
[8.3472147 ],  
[5.2013031 ],  
[8.34927612],  
[5.2953823 ],  
[9.93055323],  
[8.92738854],  
[7.61314316],  
[4.4303852 ],  
[4.89248028],  
[6.81075419],  
[8.94280325],  
[4.50713554],  
[6.05756296],  
[9.50780586],  
[6.04599748],  
[8.10591381],  
[8.16423959],  
[4.32257977],  
[9.93445953],  
[9.18365111],  
[6.53620684],  
[9.20623224],  
[7.90512826],  
[5.80455493],  
[5.26072361],  
[4.46643294],  
[8.67911213],  
[4.93670917],  
[7.0556674 ],  
[9.13340568],  
[8.23674442],  
[9.37628474],  
[7.29494673],  
[5.67395697],  
[9.90811864],  
[6.88167367],  
[4.42322404],  
[9.85790761],  
[6.05574669],  
[7.32091781],  
[5.07675957],  
[6.97295386],  
[5.58420656],  
[7.97331631],  
[9.59969769],  
[4.25142529],  
[8.43555769],


```

[9.0774013 ],
[7.60687441],
[4.56262801],
[4.19376151],
[5.46290705]])

import matplotlib.pyplot as plt

# Plotting the results
plt.scatter(X, y, color='blue', label='Actual data')
plt.plot(X, y_pred, color='red', label='Predicted line')
plt.title('Linear Regression Result')
plt.xlabel('X')
plt.ylabel('y')
plt.legend()
plt.show()

```



```

import numpy as np
import matplotlib.pyplot as plt

# Define the Linear Regression class
class LinearRegression:
    def __init__(self, learning_rate=0.01, n_iterations=1000):

```

```

        self.learning_rate = learning_rate
        self.n_iterations = n_iterations
        self.theta = None

    def fit(self, X, y):
        m = len(y)
        X_b = np.c_[np.ones((m, 1)), X] # Add bias term
        self.theta = np.random.randn(2, 1) # Random initialization
        self.mse_history = [] # Track MSE history

        for iteration in range(self.n_iterations):
            y_pred = X_b.dot(self.theta)
            mse = np.mean((y - y_pred) ** 2) # Calculate Mean Squared
Error
            self.mse_history.append(mse) # Track MSE

            gradients = 2/m * X_b.T.dot(y_pred - y)
            self.theta -= self.learning_rate * gradients

    def predict(self, X):
        X_b = np.c_[np.ones((len(X), 1)), X] # Add bias term
        return X_b.dot(self.theta)

# Generate synthetic data for testing
X = 2 * np.random.rand(100, 1)
y = 4 + 3 * X + np.random.randn(100, 1)

# Train the model
model = LinearRegression(learning_rate=0.1, n_iterations=1000)
model.fit(X, y)


# Predict values
y_pred = model.predict(X)


# Plot MSE history
plt.plot(model.mse_history)
plt.title('MSE over Iterations')
plt.xlabel('Iterations')
plt.ylabel('MSE')
plt.show()

# Plot original data and predictions
plt.scatter(X, y, color='blue', label='Actual data')
plt.plot(X, y_pred, color='red', label='Predicted line')
plt.title('Linear Regression Fit')
plt.xlabel('X')
plt.ylabel('y')
plt.legend()
plt.show()

```

MSE over Iterations





task1

```
class LinearRegression:
    def __init__(self, learning_rate=0.01, n_iterations=1000,
loss_function='mse'):
        self.learning_rate = learning_rate
        self.n_iterations = n_iterations
        self.loss_function = loss_function
        self.theta = None
        self.loss_history = []

    def compute_loss(self, y, y_pred):
        if self.loss_function == 'mse':
            return np.mean((y - y_pred) ** 2) # MSE
        elif self.loss_function == 'mae':
            return np.mean(np.abs(y - y_pred)) # MAE

    def fit(self, X, y):
        m = len(y)
        X_b = np.c_[np.ones((m, 1)), X] # Add bias term
        self.theta = np.random.randn(2, 1) # Random initialization
```

```

    for iteration in range(self.n_iterations):
        y_pred = X_b.dot(self.theta)
        loss = self.compute_loss(y, y_pred)
        self.loss_history.append(loss)

        if self.loss_function == 'mse':
            gradients = 2/m * X_b.T.dot(y_pred - y)
        elif self.loss_function == 'mae':
            gradients = X_b.T.dot(np.sign(y_pred - y)) / m

        # Ensure correct shape for gradients and theta update
        self.theta -= self.learning_rate *
gradients.reshape(self.theta.shape)

    def predict(self, X):
        X_b = np.c_[np.ones((len(X), 1)), X] # Add bias term
        return X_b.dot(self.theta)

# Generate synthetic data for testing
X = 2 * np.random.rand(100, 1)
y = 4 + 3 * X + np.random.randn(100, 1)

# Train models with both loss functions
model_mse = LinearRegression(learning_rate=0.1, n_iterations=1000,
loss_function='mse')
model_mse.fit(X, y)

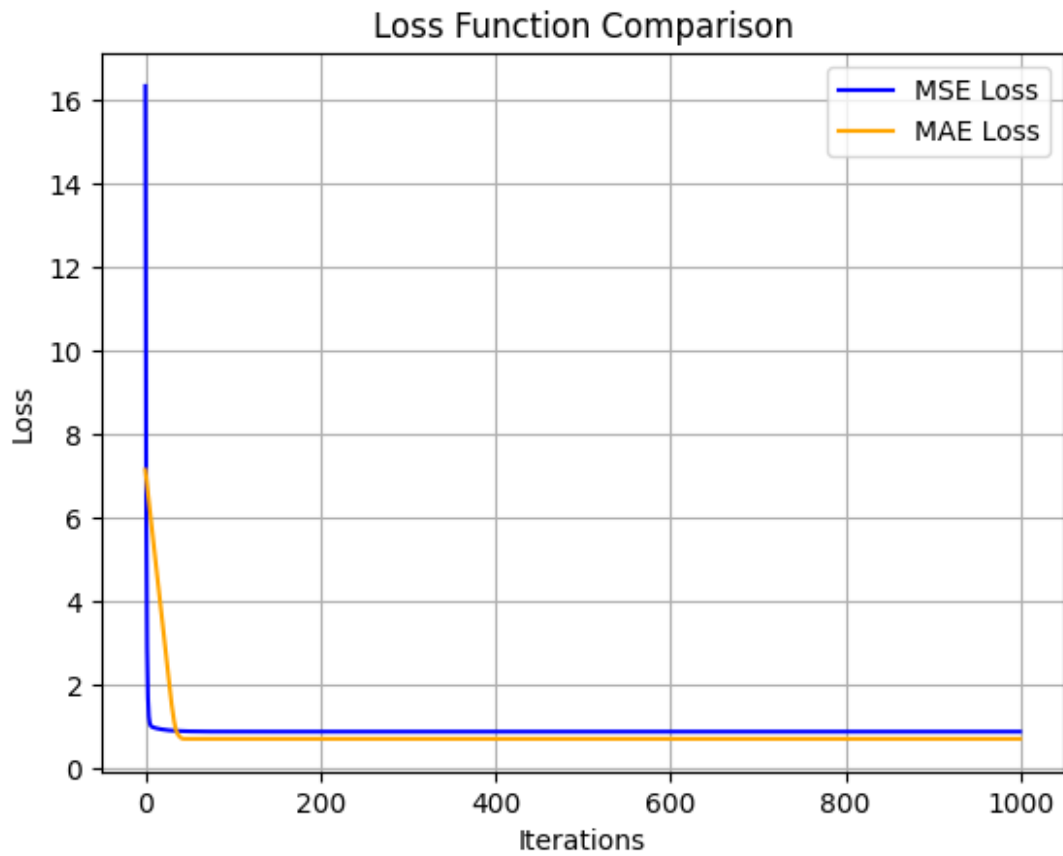
model_mae = LinearRegression(learning_rate=0.1, n_iterations=1000,
loss_function='mae')
model_mae.fit(X, y)

#compare performance

# Plotting loss history
plt.plot(model_mse.loss_history, color='blue', label='MSE Loss')
plt.plot(model_mae.loss_history, color='orange', label='MAE Loss')
plt.title('Loss Function Comparison')
plt.xlabel('Iterations')
plt.ylabel('Loss')
plt.legend()
plt.grid()
plt.show()

```





```
import numpy as np

class LinearRegression:
    def __init__(self, learning_rate=0.01, n_iterations=1000,
loss_function='mse'):
        self.learning_rate = learning_rate
        self.n_iterations = n_iterations
        self.loss_function = loss_function
        self.theta = None
        self.mse_history = []
        self.mae_history = []

    def fit(self, X, y):
        self.theta = np.random.randn(2, 1) # Random initialization

        for iteration in range(self.n_iterations):
            y_pred = self.predict(X)

            if self.loss_function == 'mse':
                loss = (1/m) * np.sum((y_pred - y) ** 2)
                self.mse_history.append(loss)
                gradients = (2/m) * X_b.T.dot(y_pred - y)
            elif self.loss_function == 'mae':
```

```

        loss = (1/m) * np.sum(np.abs(y_pred - y))
        self.mae_history.append(loss)
        gradients = (1/m) * X_b.T.dot(np.sign(y_pred - y))

        self.theta -= self.learning_rate * gradients


    def predict(self, X):
        X_b = np.c_[np.ones((len(X), 1)), X] # Add bias term
        return X_b.dot(self.theta)

# Train with MSE
model_mse = LinearRegression(learning_rate=0.1, n_iterations=1000,
loss_function='mse')
model_mse.fit(X, y)
y_pred_mse = model_mse.predict(X)

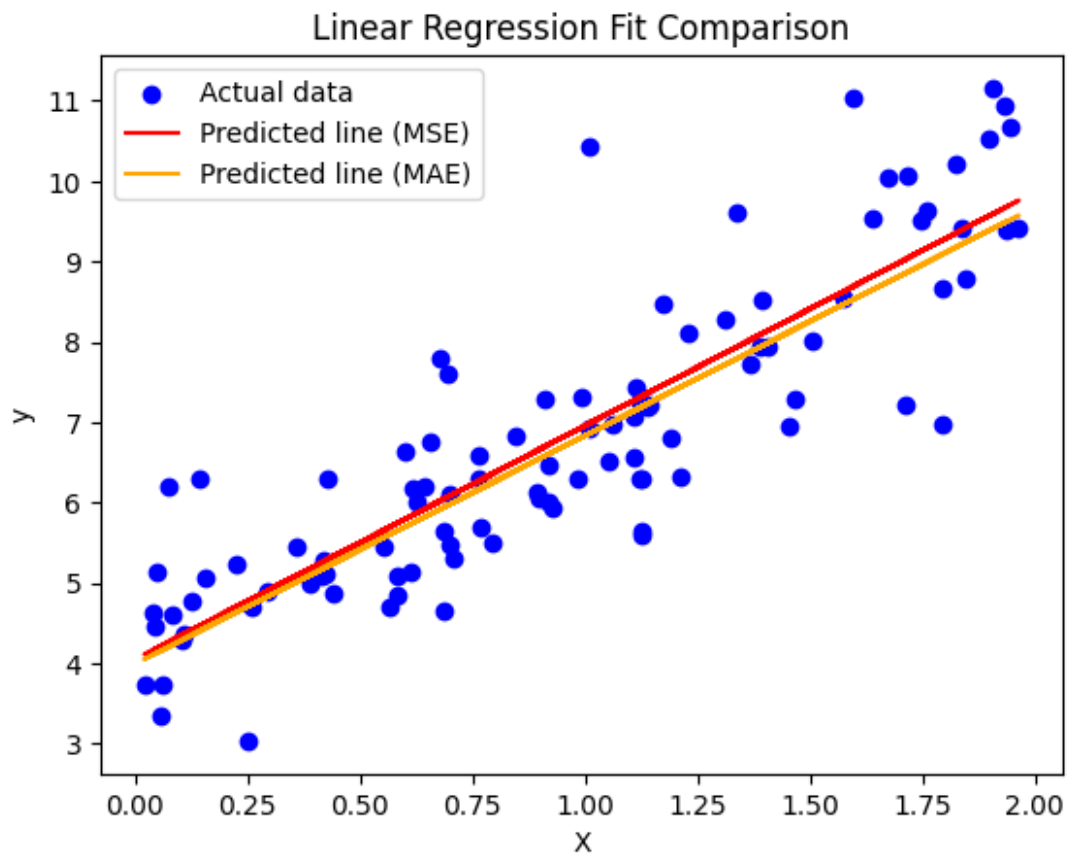
# Train with MAE
model_mae = LinearRegression(learning_rate=0.1, n_iterations=1000,
loss_function='mae')
model_mae.fit(X, y)
y_pred_mae = model_mae.predict(X)

# Plot MSE and MAE history
plt.plot(model_mse.mse_history, label='MSE', color='blue')
plt.plot(model_mae.mae_history, label='MAE', color='orange')
plt.title('Loss Function Comparison over Iterations')
plt.xlabel('Iterations')
plt.ylabel('Loss')
plt.legend()
plt.show()

```



```
# Plot original data and predictions for both models
plt.scatter(X, y, color='blue', label='Actual data')
plt.plot(X, y_pred_mse, color='red', label='Predicted line (MSE)')
plt.plot(X, y_pred_mae, color='orange', label='Predicted line (MAE)')
plt.title('Linear Regression Fit Comparison')
plt.xlabel('X')
plt.ylabel('y')
plt.legend()
plt.show()
```



```
import numpy as np  
x=np.random.randn(2, 1)
```

```
array([[2.23027983],  
       [0.9237355 ]])
```