

```

import tensorflow as tf
from tensorflow.keras.datasets import cifar10
from tensorflow.keras.utils import to_categorical

# Load dataset
(x_train_full, y_train_full), (x_test, y_test) = cifar10.load_data()

# Normalize pixel values to [0, 1]
x_train_full, x_test = x_train_full / 255.0, x_test / 255.0

# Split training data into training (60%) and validation (20%)
x_train, x_val = x_train_full[:30000], x_train_full[30000:]
y_train, y_val = y_train_full[:30000], y_train_full[30000:]

# One-hot encode labels
y_train = to_categorical(y_train)
y_val = to_categorical(y_val)
y_test = to_categorical(y_test)

Downloading data from https://www.cs.toronto.edu/~kriz/cifar-10-
python.tar.gz
170498071/170498071 ————— 4s 0us/step

import torch
import torchvision
import torchvision.transforms as transforms

# Define transforms including data augmentation
transform_train = transforms.Compose([
    transforms.RandomHorizontalFlip(),
    transforms.RandomRotation(10),
    transforms.RandomResizedCrop(32),
    transforms.ToTensor(),
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5)),
])

transform_test = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5)),
])

# Load datasets
train_dataset = torchvision.datasets.CIFAR10(root='./data',
train=True, download=True, transform=transform_train)
test_dataset = torchvision.datasets.CIFAR10(root='./data',
train=False, download=True, transform=transform_test)

# Split train dataset into training and validation
train_size = int(0.6 * len(train_dataset))
val_size = len(train_dataset) - train_size

```

```
train_dataset, val_dataset =
torch.utils.data.random_split(train_dataset, [train_size, val_size])

train_loader = torch.utils.data.DataLoader(train_dataset,
batch_size=64, shuffle=True)
val_loader = torch.utils.data.DataLoader(val_dataset, batch_size=64,
shuffle=False)
test_loader = torch.utils.data.DataLoader(test_dataset, batch_size=64,
shuffle=False)
```

Downloading <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz> to
./data/cifar-10-python.tar.gz

100%|██████████| 170M/170M [00:02<00:00, 66.5MB/s]

Extracting ./data/cifar-10-python.tar.gz to ./data
Files already downloaded and verified

```
from tensorflow.keras import layers, models, regularizers
```

```
model = models.Sequential([
    layers.Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32,
3)),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
    layers.Flatten(),
    layers.Dense(128, activation='relu'),
    layers.Dense(10, activation='softmax')
])
```

/usr/local/lib/python3.10/dist-packages/keras/src/layers/
convolutional/base_conv.py:107: UserWarning: Do not pass an
`input_shape`/`input_dim` argument to a layer. When using Sequential
models, prefer using an `Input(shape)` object as the first layer in
the model instead.

```
    super().__init__(activity_regularizer=activity_regularizer,
**kwargs)
```

```
model.compile(optimizer='adam', loss='categorical_crossentropy',
metrics=['accuracy'])
history = model.fit(x_train, y_train, epochs=20, batch_size=64,
validation_data=(x_val, y_val))
```

Epoch 1/20

469/469 ————— 52s 107ms/step - accuracy: 0.3168 - loss:
1.8542 - val_accuracy: 0.4968 - val_loss: 1.4103

Epoch 2/20

469/469 ————— 76s 95ms/step - accuracy: 0.5311 - loss:
1.3122 - val_accuracy: 0.5557 - val_loss: 1.2596

Epoch 3/20



469/469 ————— 44s 95ms/step - accuracy: 0.5896 - loss: 1.1606 - val_accuracy: 0.5820 - val_loss: 1.1960
Epoch 4/20
469/469 ————— 44s 94ms/step - accuracy: 0.6361 - loss: 1.0358 - val_accuracy: 0.6172 - val_loss: 1.1011
Epoch 5/20
469/469 ————— 82s 95ms/step - accuracy: 0.6584 - loss: 0.9734 - val_accuracy: 0.6528 - val_loss: 1.0161
Epoch 6/20
469/469 ————— 45s 97ms/step - accuracy: 0.6921 - loss: 0.8828 - val_accuracy: 0.6348 - val_loss: 1.0690
Epoch 7/20
469/469 ————— 81s 94ms/step - accuracy: 0.7103 - loss: 0.8352 - val_accuracy: 0.6557 - val_loss: 1.0030
Epoch 8/20
469/469 ————— 83s 96ms/step - accuracy: 0.7314 - loss: 0.7777 - val_accuracy: 0.6672 - val_loss: 0.9815
Epoch 9/20
469/469 ————— 82s 95ms/step - accuracy: 0.7483 - loss: 0.7261 - val_accuracy: 0.6678 - val_loss: 0.9792
Epoch 10/20
469/469 ————— 48s 101ms/step - accuracy: 0.7577 - loss: 0.6834 - val_accuracy: 0.6701 - val_loss: 1.0029
Epoch 11/20
469/469 ————— 80s 98ms/step - accuracy: 0.7835 - loss: 0.6279 - val_accuracy: 0.6752 - val_loss: 0.9863
Epoch 12/20
469/469 ————— 44s 95ms/step - accuracy: 0.8002 - loss: 0.5887 - val_accuracy: 0.6712 - val_loss: 0.9956
Epoch 13/20
469/469 ————— 82s 95ms/step - accuracy: 0.8156 - loss: 0.5387 - val_accuracy: 0.6742 - val_loss: 1.0167
Epoch 14/20
469/469 ————— 82s 95ms/step - accuracy: 0.8263 - loss: 0.5022 - val_accuracy: 0.6738 - val_loss: 1.0480
Epoch 15/20
469/469 ————— 82s 96ms/step - accuracy: 0.8400 - loss: 0.4599 - val_accuracy: 0.6716 - val_loss: 1.0632
Epoch 16/20
469/469 ————— 46s 99ms/step - accuracy: 0.8615 - loss: 0.4141 - val_accuracy: 0.6751 - val_loss: 1.1114
Epoch 17/20
469/469 ————— 80s 96ms/step - accuracy: 0.8716 - loss: 0.3811 - val_accuracy: 0.6668 - val_loss: 1.1811
Epoch 18/20
469/469 ————— 81s 94ms/step - accuracy: 0.8846 - loss: 0.3468 - val_accuracy: 0.6695 - val_loss: 1.2176
Epoch 19/20
469/469 ————— 82s 95ms/step - accuracy: 0.8958 - loss:

```
0.3159 - val_accuracy: 0.6648 - val_loss: 1.2871
Epoch 20/20
469/469 ━━━━━━━━━━━━━━━━━ 82s 95ms/step - accuracy: 0.9111 - loss:
0.2708 - val_accuracy: 0.6619 - val_loss: 1.3683
```

without regularization

```
import tensorflow as tf
from tensorflow.keras import layers, models, regularizers
from tensorflow.keras.datasets import cifar10
from tensorflow.keras.utils import to_categorical
import matplotlib.pyplot as plt
plt.figure(figsize=(10, 6))
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.title('Training and Validation Loss with Regularization')
plt.show()

# Plot training and validation accuracy
plt.figure(figsize=(10, 6))
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.title('Training and Validation Accuracy with Regularization')
plt.show()
```



Step 4: Implement Regularization

```

# Define CNN model with L2 Regularization and Dropout
def build_model():
    model = models.Sequential()

    # Convolutional Layer 1
    model.add(layers.Conv2D(32, (3, 3), activation='relu',
                             kernel_regularizer=regularizers.l2(0.001),
                             input_shape=(32, 32, 3)))
    model.add(layers.MaxPooling2D((2, 2)))

    # Convolutional Layer 2
    model.add(layers.Conv2D(64, (3, 3), activation='relu',
                             kernel_regularizer=regularizers.l2(0.001)))
    model.add(layers.MaxPooling2D((2, 2)))

    # Flatten Layer
    model.add(layers.Flatten())

    # Fully Connected Layer
    model.add(layers.Dense(128, activation='relu',
                             kernel_regularizer=regularizers.l2(0.001)))

    # Dropout Layer
    model.add(layers.Dropout(0.5))

    # Output Layer
    model.add(layers.Dense(10, activation='softmax'))

    return model

# Compile the model
model = build_model()
model.compile(optimizer='adam',
              loss='categorical_crossentropy',
              metrics=['accuracy'])

# Train the model
history = model.fit(x_train, y_train, epochs=20, batch_size=64,
                    validation_data=(x_val, y_val))

Epoch 1/20
469/469 ————— 49s 100ms/step - accuracy: 0.2612 - loss:
2.1188 - val_accuracy: 0.4669 - val_loss: 1.6078
Epoch 2/20
469/469 ————— 48s 102ms/step - accuracy: 0.4343 - loss:
1.6488 - val_accuracy: 0.5273 - val_loss: 1.4520

```

Epoch 3/20
469/469 _____ 80s 98ms/step - accuracy: 0.4887 - loss: 1.5423 - val_accuracy: 0.5511 - val_loss: 1.4020
Epoch 4/20
469/469 _____ 48s 102ms/step - accuracy: 0.5175 - loss: 1.4774 - val_accuracy: 0.5402 - val_loss: 1.4093
Epoch 5/20
469/469 _____ 80s 98ms/step - accuracy: 0.5335 - loss: 1.4364 - val_accuracy: 0.5860 - val_loss: 1.3361
Epoch 6/20
469/469 _____ 84s 102ms/step - accuracy: 0.5582 - loss: 1.3836 - val_accuracy: 0.5940 - val_loss: 1.3086
Epoch 7/20
469/469 _____ 46s 97ms/step - accuracy: 0.5701 - loss: 1.3727 - val_accuracy: 0.6128 - val_loss: 1.2776
Epoch 8/20
469/469 _____ 46s 97ms/step - accuracy: 0.5885 - loss: 1.3303 - val_accuracy: 0.6133 - val_loss: 1.2595
Epoch 9/20
469/469 _____ 82s 98ms/step - accuracy: 0.5916 - loss: 1.3094 - val_accuracy: 0.6265 - val_loss: 1.2360
Epoch 10/20
469/469 _____ 82s 99ms/step - accuracy: 0.6064 - loss: 1.2883 - val_accuracy: 0.6230 - val_loss: 1.2662
Epoch 11/20
469/469 _____ 81s 97ms/step - accuracy: 0.6117 - loss: 1.2739 - val_accuracy: 0.6240 - val_loss: 1.2493
Epoch 12/20
469/469 _____ 45s 96ms/step - accuracy: 0.6152 - loss: 1.2593 - val_accuracy: 0.6439 - val_loss: 1.2033
Epoch 13/20
469/469 _____ 83s 98ms/step - accuracy: 0.6337 - loss: 1.2204 - val_accuracy: 0.6495 - val_loss: 1.1909
Epoch 14/20
469/469 _____ 83s 100ms/step - accuracy: 0.6288 - loss: 1.2269 - val_accuracy: 0.6375 - val_loss: 1.2104
Epoch 15/20
469/469 _____ 80s 97ms/step - accuracy: 0.6428 - loss: 1.2114 - val_accuracy: 0.6582 - val_loss: 1.1724
Epoch 16/20
469/469 _____ 83s 100ms/step - accuracy: 0.6449 - loss: 1.2037 - val_accuracy: 0.6482 - val_loss: 1.2023
Epoch 17/20
469/469 _____ 46s 97ms/step - accuracy: 0.6538 - loss: 1.1890 - val_accuracy: 0.6593 - val_loss: 1.1673
Epoch 18/20
469/469 _____ 86s 106ms/step - accuracy: 0.6611 - loss: 1.1788 - val_accuracy: 0.6500 - val_loss: 1.2042
Epoch 19/20

469/469 ————— 79s 100ms/step - accuracy: 0.6668 - loss: 1.1574 - val_accuracy: 0.6626 - val_loss: 1.1650
Epoch 20/20

469/469 ————— 82s 101ms/step - accuracy: 0.6690 - loss: 1.1532 - val_accuracy: 0.6715 - val_loss: 1.1516

Evaluate the model

```
test_loss, test_accuracy = model.evaluate(x_test, y_test)
print(f"Test Loss: {test_loss}, Test Accuracy: {test_accuracy}")
```

Plot training and validation loss



```
plt.figure(figsize=(10, 6))
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.title('Training and Validation Loss with Regularization')
plt.show()
```

Plot training and validation accuracy

```
plt.figure(figsize=(10, 6))
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.title('Training and Validation Accuracy with Regularization')
plt.show()
```

313/313 ————— 5s 16ms/step - accuracy: 0.6765 - loss: 1.1458

Test Loss: 1.1531802415847778, Test Accuracy: 0.67330002784729



1. Weight Initialization: o Experiment with different weight initialization methods: □ Default initialization □ Xavier (Glorot) Initialization □ He Initialization

```
import tensorflow as tf
from tensorflow.keras import layers, models
from tensorflow.keras.datasets import cifar10
from tensorflow.keras.utils import to_categorical
import matplotlib.pyplot as plt

# Load CIFAR-10 dataset
(x_train, y_train), (x_test, y_test) = cifar10.load_data()

# Normalize the data
x_train = x_train.astype('float32') / 255.0
x_test = x_test.astype('float32') / 255.0

# One-hot encode the labels
y_train = to_categorical(y_train, 10)
y_test = to_categorical(y_test, 10)

# Split dataset into training, validation, and test sets
x_train, x_val = x_train[:30000], x_train[30000:]
y_train, y_val = y_train[:30000], y_train[30000:]

# Build the CNN Model with Variable Weight Initialization
def build_model(initializer):
    model = models.Sequential()

    # Convolutional Layer 1
    model.add(layers.Conv2D(32, (3, 3), activation='relu',
                           kernel_initializer=initializer,
                           input_shape=(32, 32, 3)))
    model.add(layers.MaxPooling2D((2, 2)))

    # Convolutional Layer 2
    model.add(layers.Conv2D(64, (3, 3), activation='relu',
                           kernel_initializer=initializer))
    model.add(layers.MaxPooling2D((2, 2)))

    # Flatten Layer
    model.add(layers.Flatten())

    # Fully Connected Layer
    model.add(layers.Dense(128, activation='relu',
                           kernel_initializer=initializer))

    # Output Layer
    model.add(layers.Dense(10, activation='softmax'))

    return model
```

```

# Experiment with different initializers
initializers = {
    'Default': 'glorot_uniform', # Default initializer in Keras
    'Xavier': tf.keras.initializers.GlorotUniform(),
    'He': tf.keras.initializers.HeNormal()
}

results = {}

for name, initializer in initializers.items():
    print(f"\nTraining model with {name} initialization...")
    model = build_model(initializer)
    model.compile(optimizer='adam',
                  loss='categorical_crossentropy',
                  metrics=['accuracy'])
    history = model.fit(x_train, y_train, epochs=5, batch_size=64,
                       validation_data=(x_val, y_val), verbose=1)
    test_loss, test_accuracy = model.evaluate(x_test, y_test)
    results[name] = {'loss': test_loss, 'accuracy': test_accuracy}

# Plot training and validation accuracy for each initializer
plt.figure(figsize=(10, 6))
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation
Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.title(f'Accuracy Curves ({name} Initialization)')
plt.legend()
plt.show()

# Compare results
for name, metrics in results.items():
    print(f"\n{name} Initialization -> Test Loss:
{metrics['loss']:.4f}, Test Accuracy: {metrics['accuracy']:.4f}")

```

Training model with Default initialization...

Epoch 1/5

469/469 ————— 48s 99ms/step - accuracy: 0.3372 - loss: 1.8161 - val_accuracy: 0.5247 - val_loss: 1.3374

Epoch 2/5

469/469 ————— 45s 96ms/step - accuracy: 0.5492 - loss: 1.2619 - val_accuracy: 0.5780 - val_loss: 1.2028

Epoch 3/5


469/469 ————— 86s 105ms/step - accuracy: 0.6154 - loss: 1.0975 - val_accuracy: 0.5977 - val_loss: 1.1627

Epoch 4/5

469/469 ————— 79s 99ms/step - accuracy: 0.6524 - loss: 0.9963 - val_accuracy: 0.6355 - val_loss: 1.0523

Epoch 5/5

```
469/469 _____ 81s 96ms/step - accuracy: 0.6845 - loss:
0.9086 - val_accuracy: 0.6508 - val_loss: 1.0149
313/313 _____ 4s 12ms/step - accuracy: 0.6469 - loss:
1.0053
```



Training model with Xavier initialization...

Epoch 1/5

```
469/469 _____ 49s 100ms/step - accuracy: 0.3326 - loss:
1.8258 - val_accuracy: 0.5082 - val_loss: 1.3674
```

Epoch 2/5

```
469/469 _____ 81s 98ms/step - accuracy: 0.5467 - loss:
1.2674 - val_accuracy: 0.5790 - val_loss: 1.2090
```

Epoch 3/5

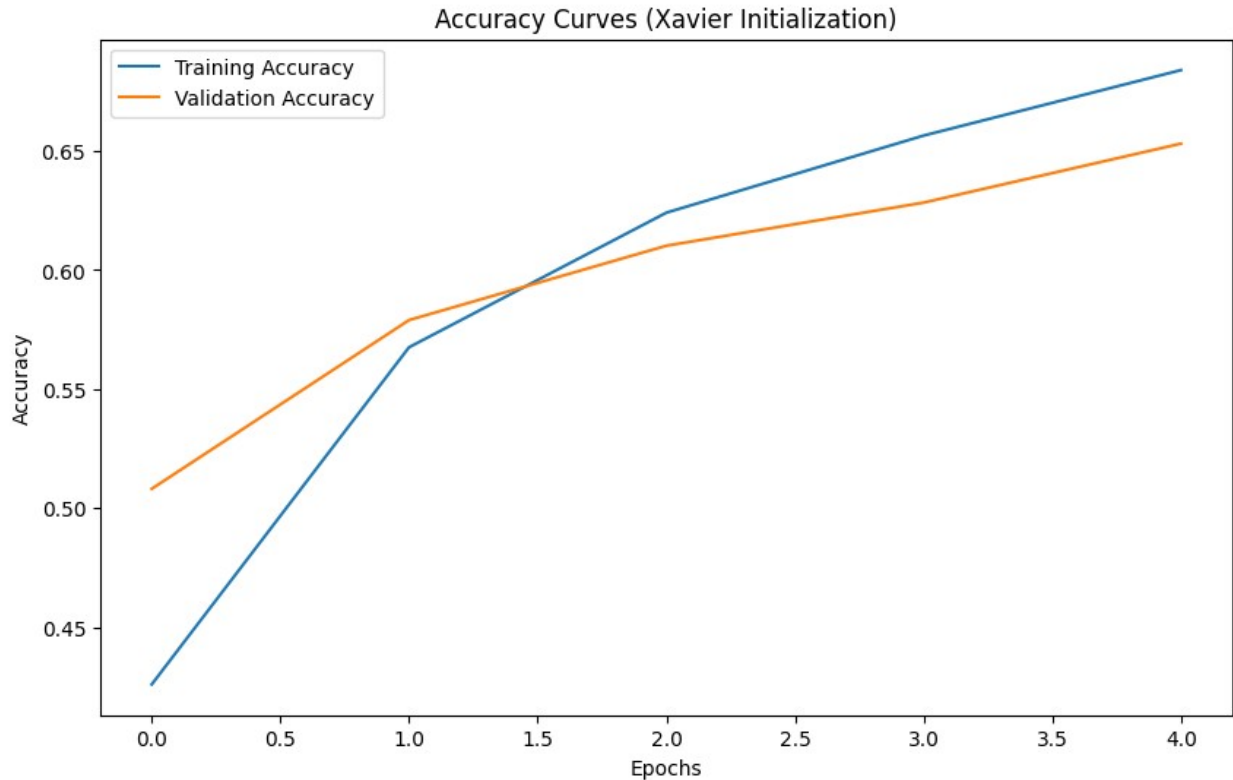
```
469/469 _____ 81s 96ms/step - accuracy: 0.6173 - loss:
1.0929 - val_accuracy: 0.6102 - val_loss: 1.1403
```

Epoch 4/5

```
469/469 _____ 84s 101ms/step - accuracy: 0.6543 - loss:
0.9896 - val_accuracy: 0.6283 - val_loss: 1.0789
```

Epoch 5/5

```
469/469 _____ 80s 96ms/step - accuracy: 0.6838 - loss:
0.9078 - val_accuracy: 0.6531 - val_loss: 1.0082
313/313 _____ 4s 13ms/step - accuracy: 0.6601 - loss:
0.9852
```



Training model with He initialization...

Epoch 1/5

469/469 ————— 51s 104ms/step - accuracy: 0.3448 - loss: 1.8173 - val_accuracy: 0.5220 - val_loss: 1.3641

Epoch 2/5

469/469 ————— 45s 97ms/step - accuracy: 0.5615 - loss: 1.2444 - val_accuracy: 0.5816 - val_loss: 1.1965

Epoch 3/5

469/469 ————— 84s 101ms/step - accuracy: 0.6268 - loss: 1.0748 - val_accuracy: 0.6189 - val_loss: 1.1007

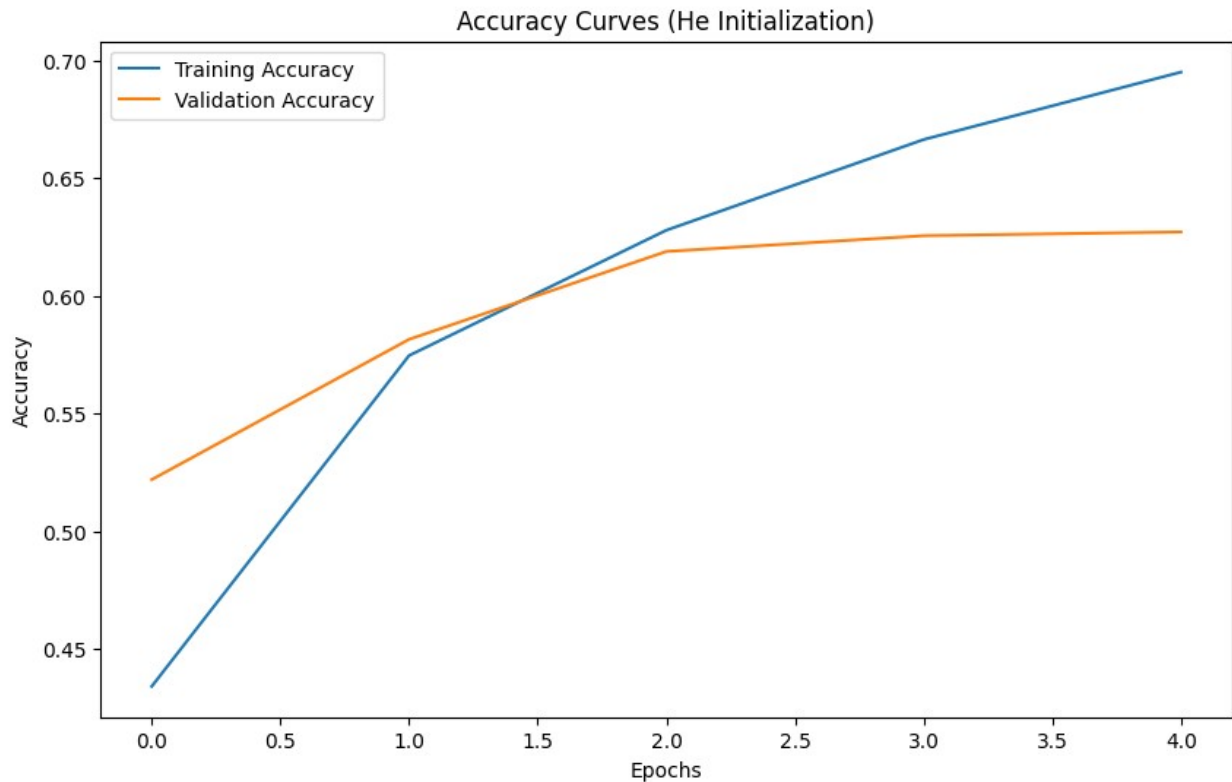
Epoch 4/5

469/469 ————— 82s 102ms/step - accuracy: 0.6667 - loss: 0.9640 - val_accuracy: 0.6256 - val_loss: 1.0776

Epoch 5/5

469/469 ————— 81s 101ms/step - accuracy: 0.7011 - loss: 0.8653 - val_accuracy: 0.6272 - val_loss: 1.1057

313/313 ————— 4s 13ms/step - accuracy: 0.6356 - loss: 1.0861



Default Initialization -> Test Loss: 1.0089, Test Accuracy: 0.6484

Xavier Initialization -> Test Loss: 0.9932, Test Accuracy: 0.6525

He Initialization -> Test Loss: 1.0925, Test Accuracy: 0.6324

Grid Search Implementation

```
from sklearn.model_selection import ParameterGrid
import tensorflow as tf

# Define hyperparameter grid
param_grid = {
    'learning_rate': [0.001, 0.01],
    'batch_size': [32, 64],
    'epochs': [10, 20],
    'optimizer': ['adam', 'sgd']
}

# Create a grid of parameters
grid = list(ParameterGrid(param_grid))

# Record results
```

```

results = []

# Loop through each parameter combination
for params in grid:
    print(f"Training with params: {params}")

    # Build model
    model = tf.keras.Sequential([
        tf.keras.layers.Conv2D(32, (3, 3), activation='relu',
input_shape=(32, 32, 3)),
        tf.keras.layers.MaxPooling2D((2, 2)),
        tf.keras.layers.Conv2D(64, (3, 3), activation='relu'),
        tf.keras.layers.MaxPooling2D((2, 2)),
        tf.keras.layers.Flatten(),
        tf.keras.layers.Dense(128, activation='relu'),
        tf.keras.layers.Dropout(0.5),
        tf.keras.layers.Dense(10, activation='softmax')
    ])

    # Compile model
    optimizer = params['optimizer']
    if optimizer == 'adam':
        optimizer =
tf.keras.optimizers.Adam(learning_rate=params['learning_rate'])
    elif optimizer == 'sgd':
        optimizer =
tf.keras.optimizers.SGD(learning_rate=params['learning_rate'])

    model.compile(optimizer=optimizer,
                  loss='categorical_crossentropy',
                  metrics=['accuracy'])

    # Train model
    history = model.fit(
        x_train, y_train,
        batch_size=params['batch_size'],
        epochs=params['epochs'],
        validation_data=(x_val, y_val),
        verbose=0
    )

    # Evaluate on validation set
    val_loss, val_accuracy = model.evaluate(x_val, y_val, verbose=0)
    results.append({'params': params, 'val_accuracy': val_accuracy})

# Find the best hyperparameter combination
best_result = max(results, key=lambda x: x['val_accuracy'])
print(f"\nBest Params: {best_result['params']}, Validation Accuracy:
{best_result['val_accuracy']:.4f}")

```

```

Training with params: {'batch_size': 32, 'epochs': 10,
'learning_rate': 0.001, 'optimizer': 'adam'}

/usr/local/lib/python3.10/dist-packages/keras/src/layers/
convolutional/base_conv.py:107: UserWarning: Do not pass an
`input_shape`/`input_dim` argument to a layer. When using Sequential
models, prefer using an `Input(shape)` object as the first layer in
the model instead.
  super().__init__(activity_regularizer=activity_regularizer,
**kwargs)

Training with params: {'batch_size': 32, 'epochs': 10,
'learning_rate': 0.001, 'optimizer': 'sgd'}
Training with params: {'batch_size': 32, 'epochs': 10,
'learning_rate': 0.01, 'optimizer': 'adam'}
Training with params: {'batch_size': 32, 'epochs': 10,
'learning_rate': 0.01, 'optimizer': 'sgd'}
Training with params: {'batch_size': 32, 'epochs': 20,
'learning_rate': 0.001, 'optimizer': 'adam'}
Training with params: {'batch_size': 32, 'epochs': 20,
'learning_rate': 0.001, 'optimizer': 'sgd'}

```

Random Search Implementation

```

import random

# Define hyperparameter space
param_space = {
    'learning_rate': [0.001, 0.01, 0.1, 0.0001],
    'batch_size': [32, 64, 128],
    'epochs': [10, 20, 30],
    'optimizer': ['adam', 'sgd', 'rmsprop']
}

# Randomly sample combinations
n_samples = 5 # Number of random combinations to test
random_combinations = [
    {k: random.choice(v) for k, v in param_space.items()} for _ in
range(n_samples)
]

results = []

for params in random_combinations:
    print(f"Training with params: {params}")

    # Model creation, compilation, and training steps remain the same
    as Grid Search.

```



```

model = tf.keras.Sequential([
    tf.keras.layers.Conv2D(32, (3, 3), activation='relu',
input_shape=(32, 32, 3)),
    tf.keras.layers.MaxPooling2D((2, 2)),
    tf.keras.layers.Conv2D(64, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D((2, 2)),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dropout(0.5),
    tf.keras.layers.Dense(10, activation='softmax')
])

optimizer = params['optimizer']
if optimizer == 'adam':
    optimizer =
tf.keras.optimizers.Adam(learning_rate=params['learning_rate'])
elif optimizer == 'sgd':
    optimizer =
tf.keras.optimizers.SGD(learning_rate=params['learning_rate'])
elif optimizer == 'rmsprop':
    optimizer =
tf.keras.optimizers.RMSprop(learning_rate=params['learning_rate'])

model.compile(optimizer=optimizer,
               loss='categorical_crossentropy',
               metrics=['accuracy'])

history = model.fit(
    x_train, y_train,
    batch_size=params['batch_size'],
    epochs=params['epochs'],
    validation_data=(x_val, y_val),
    verbose=0
)

val_loss, val_accuracy = model.evaluate(x_val, y_val, verbose=0)
results.append({'params': params, 'val_accuracy': val_accuracy})

# Find best hyperparameter combination
best_result = max(results, key=lambda x: x['val_accuracy'])
print(f"\nBest Params: {best_result['params']}, Validation Accuracy:
{best_result['val_accuracy']:.4f}")

Training with params: {'learning_rate': 0.01, 'batch_size': 64,
'epochs': 30, 'optimizer': 'adam'}

/usr/local/lib/python3.10/dist-packages/keras/src/layers/
convolutional/base_conv.py:107: UserWarning: Do not pass an
`input_shape`/`input_dim` argument to a layer. When using Sequential
models, prefer using an `Input(shape)` object as the first layer in

```

the model instead.

```
super().__init__(activity_regularizer=activity_regularizer,  
**kwargs)
```

Training with params: {'learning_rate': 0.001, 'batch_size': 128, 'epochs': 20, 'optimizer': 'rmsprop'}

Training with params: {'learning_rate': 0.0001, 'batch_size': 128, 'epochs': 20, 'optimizer': 'adam'}

Training with params: {'learning_rate': 0.001, 'batch_size': 64, 'epochs': 30, 'optimizer': 'sgd'}

Training with params: {'learning_rate': 0.01, 'batch_size': 32, 'epochs': 10, 'optimizer': 'sgd'}

Best Params: {'learning_rate': 0.001, 'batch_size': 128, 'epochs': 20, 'optimizer': 'rmsprop'}, Validation Accuracy: 0.6619

1. Observe Gradient Norms and Gradient Clipping

```
import tensorflow as tf  
import numpy as np  
import matplotlib.pyplot as plt  
from sklearn.metrics import classification_report, confusion_matrix  
import seaborn as sns  
from tensorflow.keras.datasets import cifar10  
from tensorflow.keras.utils import to_categorical  
  
# Load CIFAR-10 dataset  
(x_train_full, y_train_full), (x_test, y_test) = cifar10.load_data()  
  
# Normalize pixel values to [0, 1]  
x_train_full, x_test = x_train_full / 255.0, x_test / 255.0  
  
# Split training data into training (60%) and validation (20%)  
x_train, x_val = x_train_full[:30000], x_train_full[30000:]  
y_train, y_val = y_train_full[:30000], y_train_full[30000:]  
  
# One-hot encode labels  
y_train = to_categorical(y_train)  
y_val = to_categorical(y_val)  
y_test = to_categorical(y_test)  
  
# Define the model  
model = tf.keras.Sequential([  
    tf.keras.layers.Conv2D(32, (3, 3), activation='relu',  
input_shape=(32, 32, 3)),  
    tf.keras.layers.MaxPooling2D((2, 2)),  
    tf.keras.layers.Conv2D(64, (3, 3), activation='relu'),  
    tf.keras.layers.MaxPooling2D((2, 2)),  
    tf.keras.layers.Flatten(),  
    tf.keras.layers.Dense(128, activation='relu'),
```

```

        tf.keras.layers.Dense(10, activation='softmax')
    ])

# Compile the model
model.compile(optimizer=tf.keras.optimizers.Adam(),
              loss='categorical_crossentropy',
              metrics=['accuracy'])

# Build the model (necessary for gradient monitoring)
model.build(input_shape=(None, 32, 32, 3))

# Gradient Monitoring Callback
class GradientMonitor(tf.keras.callbacks.Callback):
    def __init__(self):
        self.gradient_norms = []

    def on_train_batch_end(self, batch, logs=None):
        with tf.GradientTape() as tape:
            # Forward pass using dummy data to avoid undefined input error
            dummy_input = tf.random.uniform((32, 32, 32, 3)) # Batch size of 32
            logits = self.model(dummy_input, training=True)
            dummy_target = tf.one_hot(tf.random.uniform((32,)),
maxval=10, dtype=tf.int32), 10)
            loss = self.model.compiled_loss(dummy_target, logits,
regularization_losses=self.model.losses)

            # Compute gradients
            gradients = tape.gradient(loss,
self.model.trainable_variables)
            # Calculate norm for each layer
            norms = [tf.norm(g).numpy() if g is not None else 0 for g in
gradients]
            self.gradient_norms.append(sum(norms))

# Apply Gradient Clipping
optimizer = tf.keras.optimizers.Adam(clipnorm=1.0) # Clip gradients by norm
model.compile(optimizer=optimizer, loss='categorical_crossentropy',
metrics=['accuracy'])

# Gradient monitor instance
gradient_monitor = GradientMonitor()

# Train the model
history = model.fit(
    x_train, y_train,
    epochs=5,
    batch_size=32,

```

```

        validation_data=(x_val, y_val),
        callbacks=[gradient_monitor]
    )

    # Plot gradient norms
    plt.plot(gradient_monitor.gradient_norms)
    plt.title("Gradient Norms During Training")
    plt.xlabel("Training Batch")
    plt.ylabel("Gradient Norm")
    plt.show()

    # Evaluate the model
    test_loss, test_accuracy = model.evaluate(x_test, y_test, verbose=0)
    print(f"Test Loss: {test_loss:.4f}, Test Accuracy:
    {test_accuracy:.4f}")

    # Generate Confusion Matrix and Classification Report
    y_pred = model.predict(x_test)
    y_pred_classes = np.argmax(y_pred, axis=1)
    y_true_classes = np.argmax(y_test, axis=1)

    conf_matrix = confusion_matrix(y_true_classes, y_pred_classes)
    plt.figure(figsize=(10, 8))
    sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues',
    xticklabels=list(range(10)), yticklabels=list(range(10)))
    plt.xlabel('Predicted Class')
    plt.ylabel('True Class')
    plt.title('Confusion Matrix')
    plt.show()

    # Classification Report
    class_report = classification_report(y_true_classes, y_pred_classes,
    target_names=[str(i) for i in range(10)])
    print("Classification Report:\n", class_report)

    Downloading data from https://www.cs.toronto.edu/~kriz/cifar-10-
    python.tar.gz
    170498071/170498071 ————— 3s 0us/step

    /usr/local/lib/python3.10/dist-packages/keras/src/layers/
    convolutional/base_conv.py:107: UserWarning: Do not pass an
    `input_shape`/`input_dim` argument to a layer. When using Sequential
    models, prefer using an `Input(shape)` object as the first layer in
    the model instead.
      super().__init__(activity_regularizer=activity_regularizer,
      **kwargs)

```

Epoch 1/5


```

/usr/local/lib/python3.10/dist-packages/keras/src/backend/tensorflow/
trainer.py:617: UserWarning: `model.compiled_loss()` is deprecated.


```

```
Instead, use `model.compute_loss(x, y, y_pred, sample_weight,
training)`.
warnings.warn(
```

```
938/938 _____ 137s 141ms/step - accuracy: 0.3709 -
loss: 1.7299 - val_accuracy: 0.4926 - val_loss: 1.3959
Epoch 2/5
938/938 _____ 124s 132ms/step - accuracy: 0.5794 -
loss: 1.1774 - val_accuracy: 0.6053 - val_loss: 1.1303
Epoch 3/5
938/938 _____ 140s 130ms/step - accuracy: 0.6469 -
loss: 1.0121 - val_accuracy: 0.6345 - val_loss: 1.0451
Epoch 4/5
938/938 _____ 147s 136ms/step - accuracy: 0.6826 -
loss: 0.9048 - val_accuracy: 0.6492 - val_loss: 1.0210
Epoch 5/5
938/938 _____ 136s 130ms/step - accuracy: 0.7251 -
loss: 0.7895 - val_accuracy: 0.6683 - val_loss: 0.9717
```



```
Test Loss: 0.9714, Test Accuracy: 0.6649
313/313 _____ 6s 18ms/step
```



Classification Report:

	precision	recall	f1-score	support
0	0.66	0.76	0.71	1000
1	0.73	0.84	0.78	1000
2	0.64	0.43	0.51	1000
3	0.53	0.47	0.50	1000
4	0.65	0.52	0.58	1000
5	0.51	0.64	0.57	1000
6	0.64	0.82	0.72	1000
7	0.75	0.69	0.72	1000
8	0.82	0.72	0.77	1000
9	0.75	0.76	0.75	1000
accuracy			0.66	10000
macro avg	0.67	0.66	0.66	10000

weighted avg	0.67	0.66	0.66	10000
--------------	------	------	------	-------